CS61C

**Great Ideas
in
Computer Architecture**
(a.k.a. Machine Structures)

UC Berkeley
Teaching Professor
Dan Garcia

UC Berkeley
Professor
Bora Nikolić

Pipelining

cs61c.org

New-School Machine Structures

Software

Parallel Requests
Assigned to computer
    e.g., Search "Cats"

Parallel Threads
Assigned to core e.g., Lookup, Ads

Parallel Instructions
    >1 instruction @ one time
    e.g., 5 pipelined instructions

Parallel Data
>1 data item @ one time
    e.g., Add of 4 pairs of words

Hardware descriptions
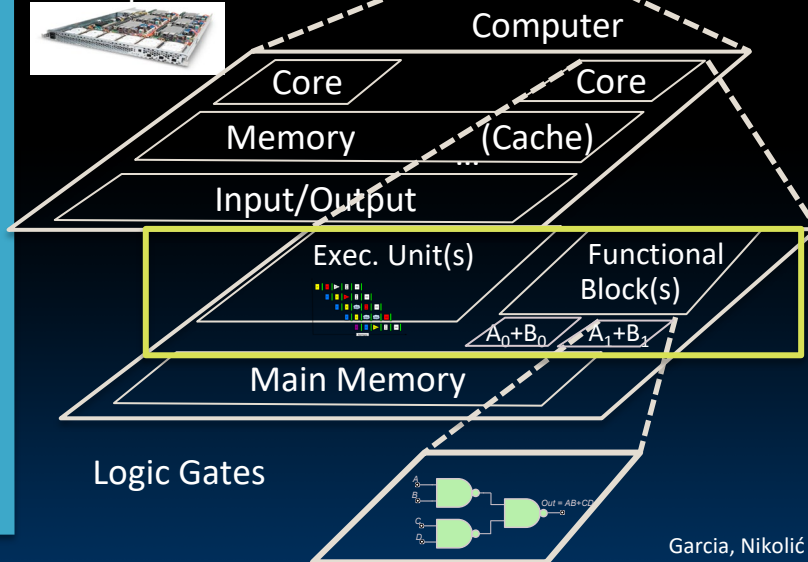All gates work in parallel at same time

Harness Parallelism & Achieve High Performance
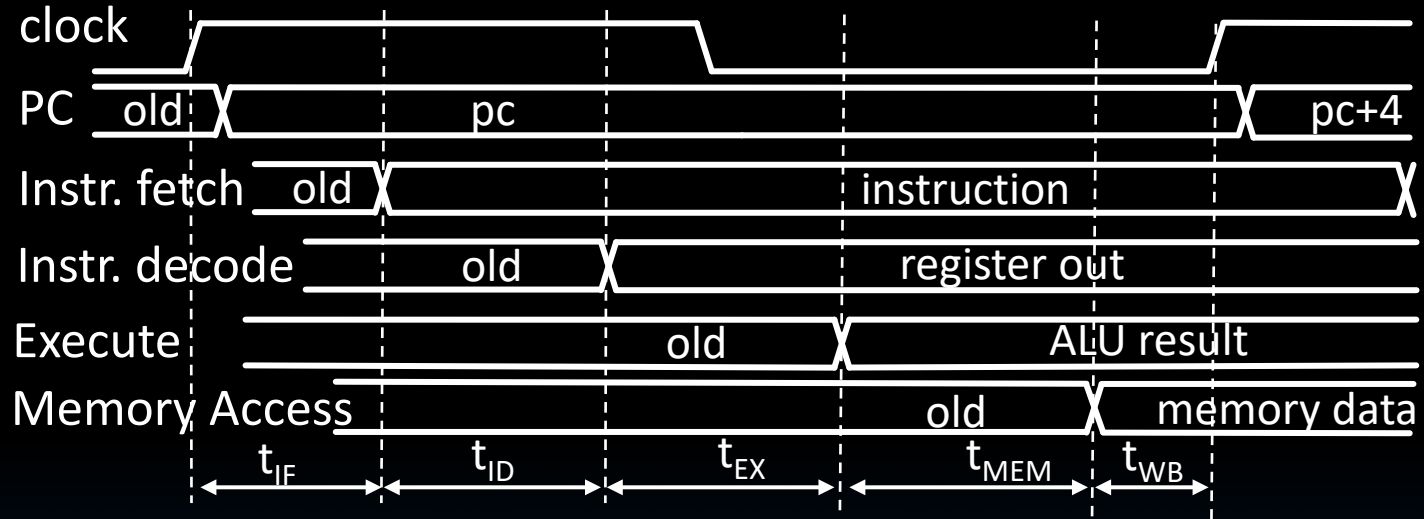
Hardware

Warehouse Scale Computer

Smart Phone

Computer

Core   Core

Memory   (Cache)

Input/Output

Exec. Unit(s)   Functional Block(s)

$A_0 + B_0$ $A_1 + B_1$

Main Memory

Logic Gates

$Out = AB + CC$

Berkeley
UNIVERSITY OF CALIFORNIA

	IF	ID	EX	MEM	WB	Total
	I-MEM	Reg Read	ALU	D-MEM	Reg W	
	200 ps	100 ps	200 ps	200 ps	100 ps	**800 ps**

Instruction Timing

Instr	IF = 200ps	ID = 100ps	ALU = 200ps	MEM=200ps	WB = 100ps	Total
add	X	X	X		X	600ps
beq	X	X	X			500ps
jal	X	X	X		X	600ps
lw	X	X	X	X	X	800ps
sw	X	X	X	X		700ps

- Maximum clock frequency
  - f_{max} = 1/800ps = 1.25 GHz

- "Our" Single-cycle RISC-V CPU executes instructions at 1.25 GHz
  - 1 instruction every 800 ps
- Can we improve its performance?
  - What do we mean with this statement?
  - Not so obvious:
    - Quicker response time, so one job finishes faster?
    - More jobs per unit time (e.g. web server returning pages, spoken words recognized)?
    - Longer battery life?

Garcia, Nikolić

Transportation Analogy

	Sports Car	Bus
Passenger Capacity	2	50
Travel Speed	200 mph	50 mph
Gas Mileage	5 mpg	2 mpg

50 Mile trip (assume they return instantaneously)

	Sports Car	Bus
Travel Time	15 min	60 min
Time for 100 passengers	750 min (50 2-person trips)	120 min (two 50-person trips)
Gallons per passenger	5 gallons	0.5 gallons

Garcia, Nikolić

Berkeley
UNIVERSITY OF CALIFORNIA

Computer Analogy

Transportation	Computer
Trip Time	Program execution time: e.g. time to update display
Time for 100 passengers	Throughput: e.g. number of server requests handled per hour
Gallons per passenger	Energy per task*: e.g. how many movies you can watch per battery charge or energy bill for datacenter

* <u>Note</u>:   Power is not a good measure, since low-power CPU might run for a long time to complete one task consuming more energy than faster computer running at higher power for a shorter time

Garcia, Nikolić

Processor Performance Iron Law

$$\frac{\text{Time}}{\text{Program}} = \frac{\text{Instructions}}{\text{Program}} * \frac{\text{Cycles}}{\text{Instruction}} * \frac{\text{Time}}{\text{Cycle}}$$

CPI = <u>C</u>ycles <u>P</u>er <u>I</u>nstruction

Determined by

- Task

- Algorithm, e.g. $O(N^2)$ vs $O(N)$

- Programming language

- Compiler

- Instruction Set Architecture (ISA)

Berkeley
UNIVERSITY OF CALIFORNIA

Determined by

- ISA

- Processor implementation (or *microarchitecture*)

- E.g. for "our" single-cycle RISC-V design, CPI = 1

- Complex instructions (e.g. `strcpy`), CPI >> 1

- Superscalar processors, CPI < 1 (next lectures)

Determined by

- Processor microarchitecture (determines critical path through logic gates)

- Technology (e.g. 5nm versus 28nm)

- Power budget (lower voltages reduce transistor speed)

Garcia, Nikolić

- For some task (e.g. image compression) …

	Processor A	Processor B
# Instructions	1 Million	1.5 Million
Average CPI	2.5	1
Clock rate f	2.5 GHz	2 GHz
Execution time	1 ms	0.75 ms

Processor B is faster for this task, despite executing more instructions and having a slower clock rate!

Garcia, Nikolić

Energy Efficiency

Where Does Energy Go in CMOS?

Symbol (INV)

Schematic

V_{DD}

A

Out

A

Out

Leakage

Charging capacitors (CV^2)

(30%)

(70%)

Garcia, Nikolić

Berkeley
UNIVERSITY OF CALIFORNIA

$$\frac{\text{Energy}}{\text{Program}} = \frac{\text{Instructions}}{\text{Program}} * \frac{\text{Energy}}{\text{Instruction}}$$

$$\frac{\text{Energy}}{\text{Program}} \propto \frac{\text{Instructions}}{\text{Program}} * C \, V^2$$

"Capacitance" depends on technology,
processor features
e.g. # of cores

Supply voltage,
e.g. 1V

Want to reduce capacitance and voltage to reduce energy/task

Berkeley
UNIVERSITY OF CALIFORNIA

Garcia, Nikolić

- **"Next-generation" processor**
  - C (Moore's Law):   -15 %
  - Supply voltage, V_{sup}:   -15 %
  - Energy consumption: $0 - (1-0.85^3) =$ **-39 %**

- **Significantly improved energy efficiency thanks to**
  - Moore's Law **AND**
  - Reduced supply voltage

48 Years of Microprocessor Trend Data

Transistors (thousands)

Single-Thread Performance (SpecINT x 10^3)

Frequency (MHz)

Typical Power (Watts)

Number of Logical Cores

Original data up to the year 2010 collected and plotted by M. Horowitz, F. Labonte, O. Shacham, K. Olukotun, L. Hammond, and C. Batten
New plot and data collected for 2010-2019 by K. Rupp

Garcia, Nikolić

- In recent years, industry has not been able to reduce supply voltage much, as reducing it further would mean increasing "leakage power" where transistor switches don't fully turn off (more like dimmer switch than on-off switch)

- Also, size of transistors and hence capacitance, not shrinking as much as before between transistor generations

  - Need to go to 3D

- Power becomes a growing concern – the "power wall"

Garcia, Nikolić

Performance = Power * Energy Efficiency

(Tasks/Second) (Joules/Second) (Tasks/Joule)

- Energy efficiency (e.g., instructions/Joule) is key metric in all computing devices
- For power-constrained systems (e.g., 20MW datacenter), need better energy efficiency to get more performance at same power
- For energy-constrained systems (e.g., 1W phone), need better energy efficiency to prolong battery life
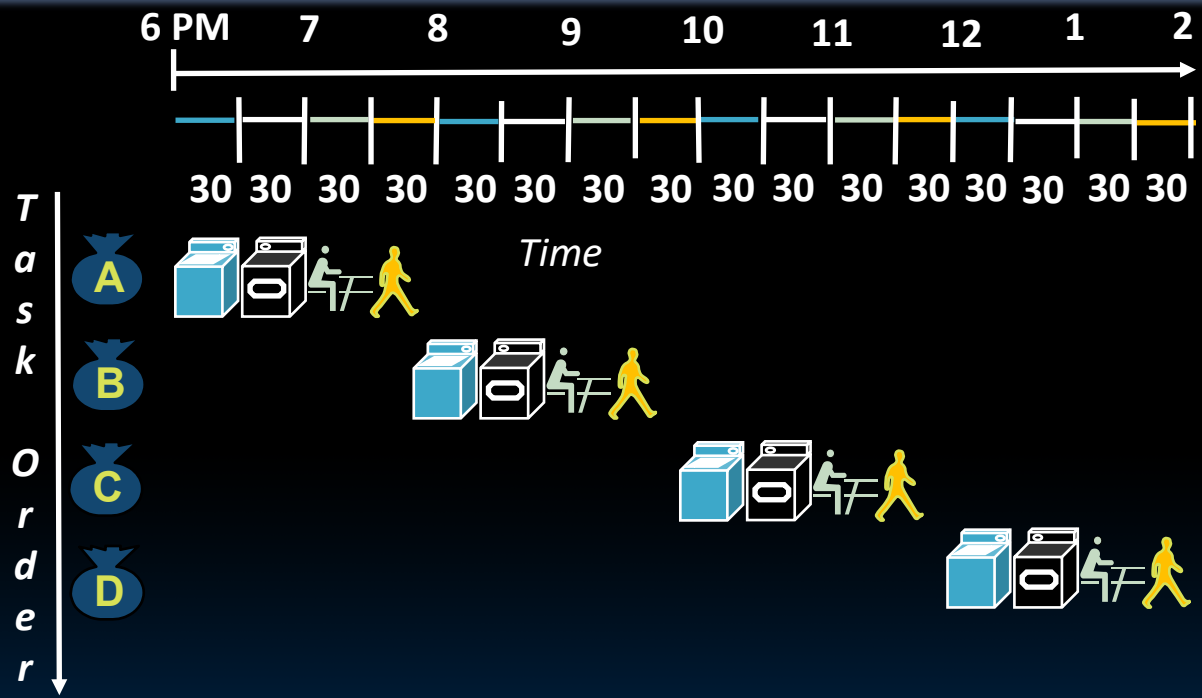
Garcia, Nikolić

- Avi, Bora, Caroline, Dan each have one load of clothes to wash, dry, fold, and put away
  - Washer takes 30 minutes

  - Dryer takes 30 minutes

  - "Folder" takes 30 minutes

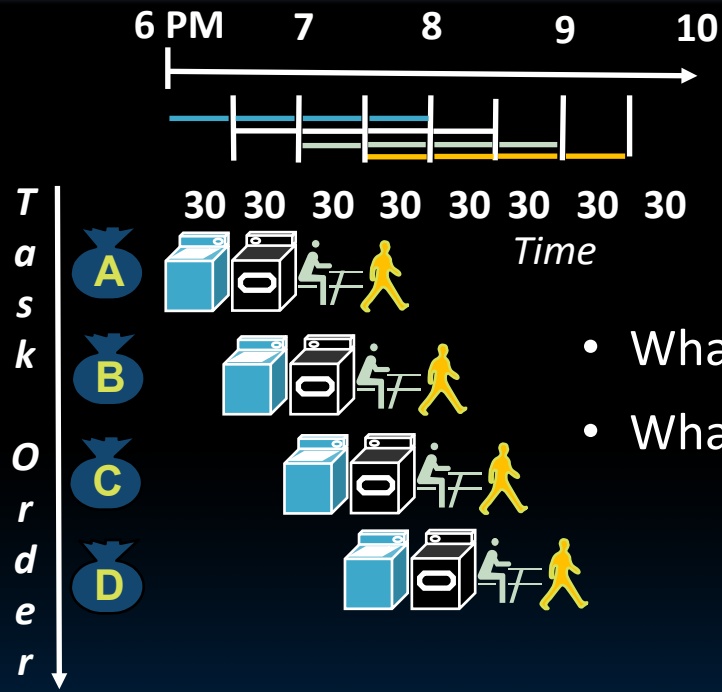 - "Stasher" takes 30 minutes to put clothes into drawers
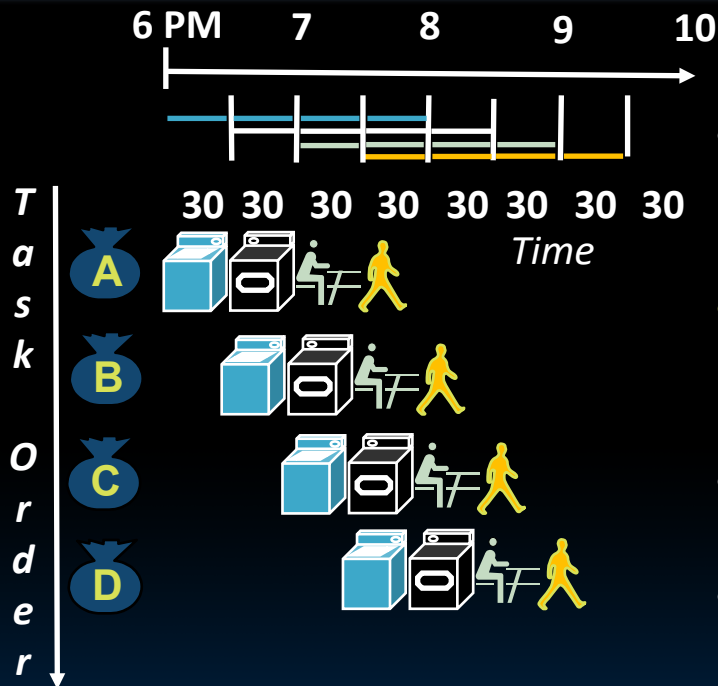
Sequential laundry takes 8 hours for 4 loads!

Garcia, Nikolić

6 PM   7   8   9   10

30 30 30 30 30 30 30 30

Time

T
a
s
k

A

B

O
r
d
e
r

C

D

- What happens *sequentially*?

- What happens *simultaneously*?

Pipelined laundry takes 3.5 hours for 4 loads!

Berkeley
UNIVERSITY OF CALIFORNIA

6 PM   7   8   9   10

Time

T a s k   O r d e r

30 30 30 30 30 30 30 30

A

B

C

D

- Pipelining doesn't help latency of single task, it helps throughput of entire workload
- Multiple tasks operating simultaneously using different resources
- Potential speedup = Number of pipe stages
- Time to "fill" pipeline and time to "drain" it reduces speedup: 2.3X v. 4X in this example

Berkeley
UNIVERSITY OF CALIFORNIA
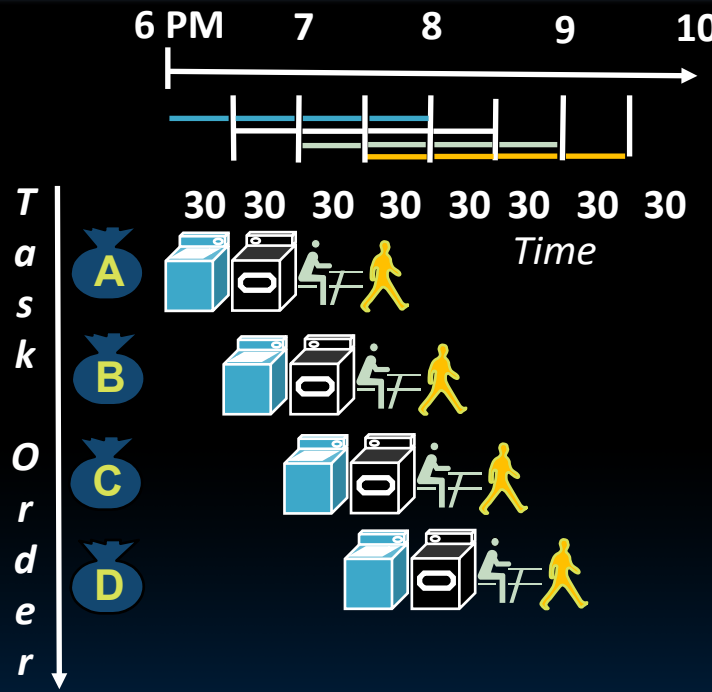
Suppose:
- new Washer takes 20 minutes
- new Stasher takes 20 minutes.

How much faster is pipeline?

Pipeline rate limited by **slowest** pipeline stage

Unbalanced lengths of pipe stages reduce speedup